

Benchmarking: Deequ **vs.** DuckDQ's Performance on Static and Streaming Data

CMPT 984 Project Presentation

Xin Tan & Yuyu Lai

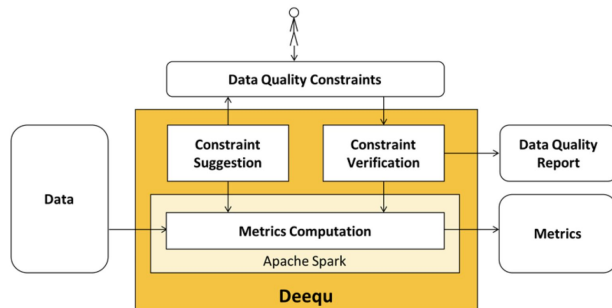
Apr 11, 2024



Problem & Motivation

Automating Data Quality Verification: Deequ

Library built on top of Apache Spark for defining "unit tests for data".



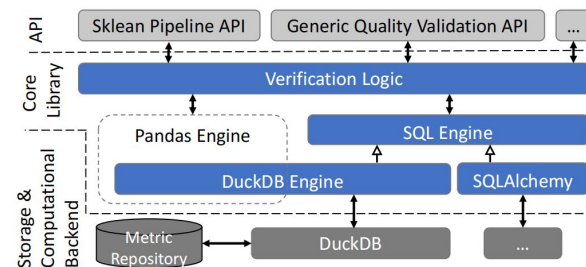
Open Source: <https://github.com/aws-labs/deequ>

- Metrics Computation
- Constraint Suggestion
- Constraint Verification
- Metrics Repository

[1] Sebastian Schelter, Dustin Lange, Philipp Schmidt, Meltem Celikel, Felix Biessmann, and Andreas Grafberger. 2018. Automating large-scale data quality verification.

Embeddable Data Quality Validation: DuckDQ

Python library that provides a fluent API for data quality checks.



Open Source: <https://github.com/tdoehmen/duckdq>

- Inspired by Deequ
- Excels at small to medium sized datasets
- Lightweight
- Outperforms existing solutions in runtime

[2] Till Doehmen, Mark Raasveldt, Hannes Muehleisen, and Sebastian Schelter. 2021. DuckDQ: Data Quality Assertions for Machine Learning Pipelines

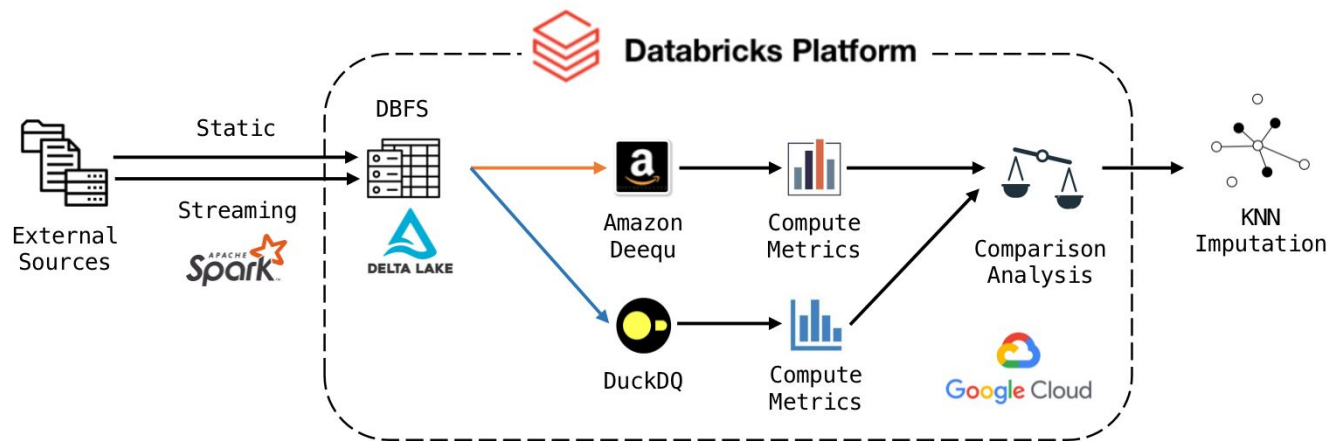
Bit Flip

Previous Work:

- ❖ Deequ has only been evaluated on fixed data.
- ❖ No comparison with other state-of-the-art systems have been made.
- ❖ Lack of data imputation feature.

Our Contribution:

- ❖ Evaluate Deequ's performance on streaming data.
- ❖ Benchmark Deequ against DuckDQ.
- ❖ Based on the suggestions by Deequ, perform data imputation.



Implementation Methods

Dataset: 'stockTicks.json' (1000 rows, 153KB)

```
{"symbol": "BLDR", "date": "08/01/2019", "time": "6:35:51", "price": 3.5343, "quantity": 3029.4181, "buysell": "sell", "ordertype": "batch", "ipaddr": "167.5.227.249"}
```

Platform: Databricks notebook workspace, Google Cloud Platform



Google Cloud




Databricks

Table 1. Experiment Setup

Compute Resources	n2-highmem-4 (1 Driver, 32 GB Memory, 4 Cores)
Runtime Version	13.3 LTS (includes Apache Spark 3.4.1, Scala 2.12)
Static Data	Uploaded file to DBFS
Streaming Data	Read and store as parquet, simulate streaming, write to Delta Tables

Streaming Data

Simulate Streaming: analysis in a streaming context without real-time data generation.



1. **Read and Store as Parquet:** We use Spark to read and repartition the JSON data into 100 partitions, then write to a temporary folder in Parquet format.
2. **Streaming:** simulate a data stream from the parquet files using `'readStream'` with `.option("maxFilesPerTrigger",1)` to limit file processing per trigger.
3. **Write to Delta Tables:** Stream-read data is written to 'trades_delta' Delta table for analysis.

Streaming Data

Cmd 15

```
1 %scala
2 // parse the schema for the source parquet
3 val schema = base_df.schema
4
5 // start the stream
6 spark.readStream
7   .schema(schema)
8   .format("parquet")
9   .option("maxFilesPerTrigger",1)
10  .load(data_path)
11  .writeStream.format("delta")
12  .option("failOnDataLoss", false)
13  .option("checkpointLocation", checkpoint_path)
14  .format("delta").table("trades_delta")
```

▼ (1) Spark Jobs

▶ Job 1794 [View](#) (Stages: 2/2)

▶ 87af1d53-422e-4889-b5c2-fdefed208ada *Last updated: 3 days ago*

```
schema: org.apache.spark.sql.types.StructType = StructType(StructField(buysell,StringType,true),StructField(date,StringType,true),StructField(ipaddr,StringType,true),StructField(ordertype,StringType,true),StructField(price,DoubleType,true),StructField(quantity,DoubleType,true),StructField(symbol,StringType,true),StructField(time,StringType,true))
res7: org.apache.spark.sql.streaming.StreamingQuery = org.apache.spark.sql.execution.streaming.StreamingQueryWrapper@24df4553
```


Command complete

Dashboard

Raw Data


Input vs. Processing Rate
records per second

1 rec/s 0.5 rec/s
Input rate Processing rate

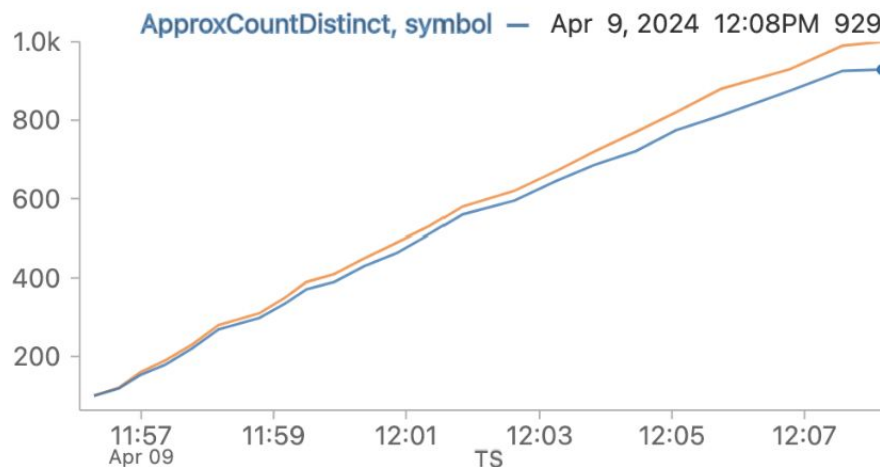


Batch Duration
in seconds

6.4 s 21.1 s
Average Latest



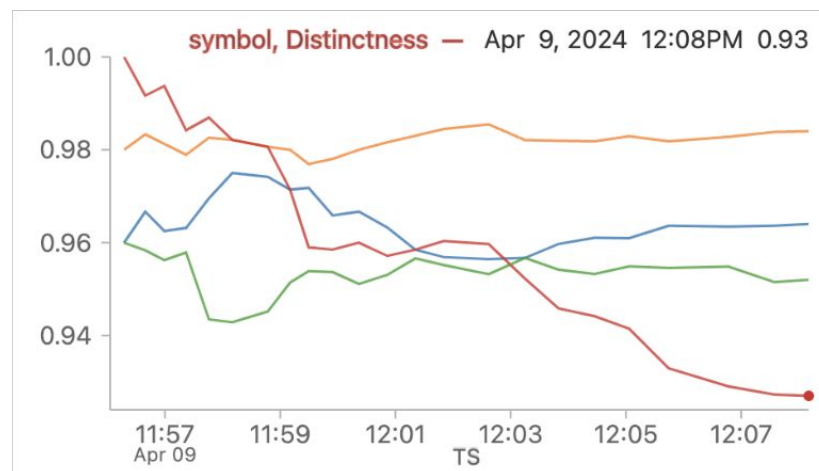
Historical Distinctness



name, instance

■ ApproxCountDistinct,
■ Size, *

Historical Completeness




instance, name

■ quantity, Completeness ■ symbol, Distinctness
■ ipaddr, Completeness ■ price, Completeness

Evaluation 1: Streaming vs. Static Results


Streaming Results



Constraint

■ CompletenessConstraint(Completeness(ipaddr))
■ CompletenessConstraint(Completeness(quantity))

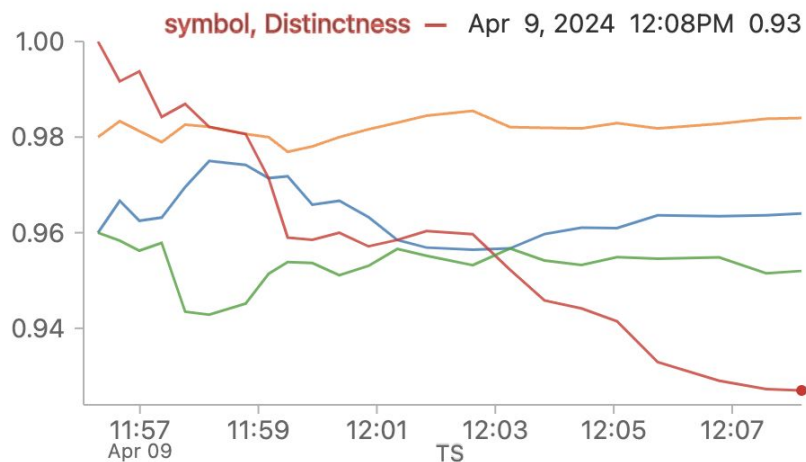
Static Results



■ CompletenessConstraint(Completeness(price))
■ DistinctnessConstraint(Distinctness(symbol))

Evaluation 2: Incremental Computation vs. Batch

Deequ



Constraint


CompletenessConstraint(Completeness(ipaddr))

CompletenessConstraint(Completeness(quantity))

CompletenessConstraint(Completeness(price))

DistinctnessConstraint(Distinctness(symbol))

DuckDQ:



Evaluation 3: Performance over Static Data

Table 2. Runtime in Seconds

Tools	Completeness	Uniqueness	Distinctness	isApproxQuantile	NonNegative	Full Run
Deequ	0.49	3.6	3.34	0.68	1.03	5.89
DuckDQ	0.16	0.21	0.21	0.15	0.18	0.22

Table 3. Verification Results

Constraint	DuckDQ	Deequ
DistinctnessConstraint(Distinctness(buysell))	0.003	0.002
UniquenessConstraint(Uniqueness(ipaddr))	0.957	1
UniquenessConstraint(Uniqueness(price))	0.966	0.983
UniquenessConstraint(Uniqueness(quantity))	0.909	0.952

KNN Imputation

To apply the findings of Deequ's data verification check:

- ❖ Experimented with KNN imputation.
- ❖ KNN is shown to outperform other methods on numeric datasets.

```
import pandas as pd
from sklearn.impute import KNNImputer
from sklearn.preprocessing import OneHotEncoder

numeric_columns = ['price', 'quantity']
numeric_df = df[numeric_columns]

# Perform KNN Imputation
imputer = KNNImputer(n_neighbors=5)
numeric_imputed = imputer.fit_transform(numeric_df)
```

Dataset with missing values:

	price	quantity
0	NaN	3338.6100
53	NaN	1667.2239
66	16.2680	NaN
68	NaN	588.3304
81	23.7744	NaN
...
909	21.8370	NaN
920	21.3060	NaN
923	27.3420	NaN
926	19.5546	NaN
940	18.4585	NaN

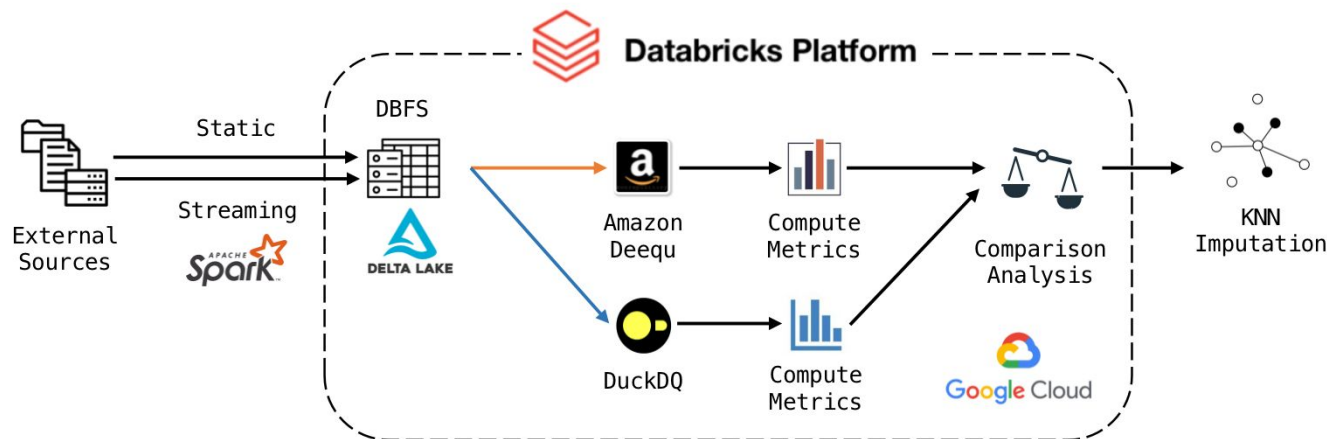


After imputation:

	price	quantity
0	26.33890	3338.61000
53	30.30720	1667.22390
66	16.26800	2129.36564
68	32.19774	588.33040
81	23.77440	2055.06980
...
909	21.83700	2054.16406
920	21.30600	2865.91608
923	27.34200	1317.02882
926	19.55460	1738.44426
940	18.45850	2572.54356

Conclusion

1. **Capacity:** Both DeeQu and DuckDQ has the ability to handle multiple types of data.
2. **For static and small to medium size data:** DuckDQ is better due to its runtime.
3. **For streaming data:** DeeQu is better because it incorporates **incremental computation** so that the overall metrics of streaming data could be observed.

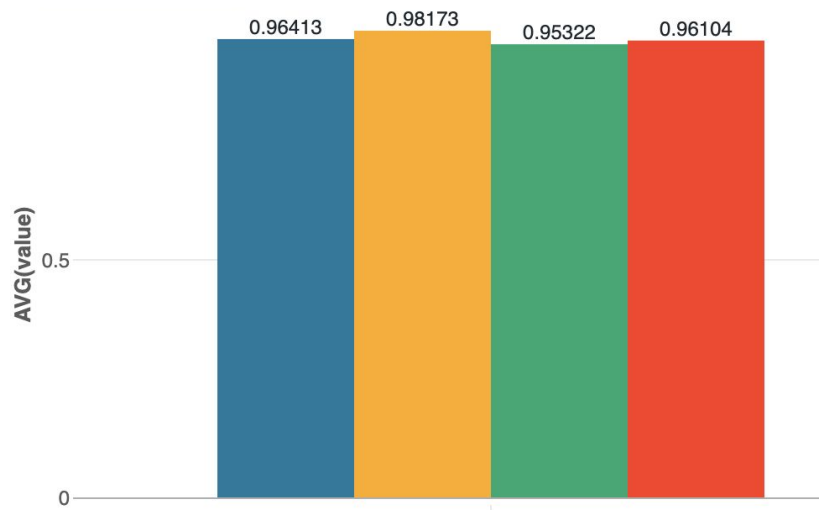


Thank you!


Q & A

Appendix - 1 Streaming Results

Deequ:



DuckDQ:

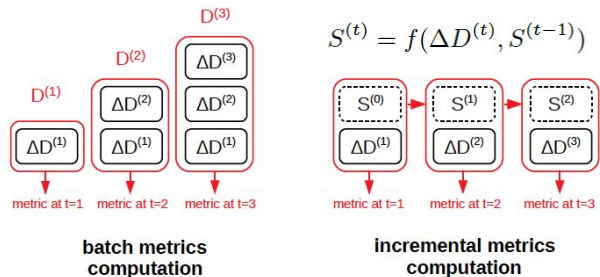


Constraint

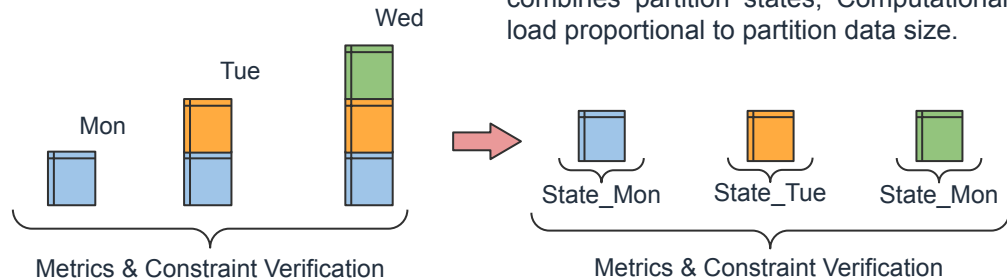
- CompletenessConstraint(Completeness(ipaddr))
- CompletenessConstraint(Completeness(price))
- CompletenessConstraint(Completeness(quantity))
- DistinctnessConstraint(Distinctness(symbol))

Appendix - 2 Incremental Computation of Metrics for Growing Datasets

Instead of repeatedly running the batch computation on growing input data D, running an incremental computation that only needs to consume the latest dataset delta $\Delta D(t)$ and a state S of the computation.



Example: logs with daily partitions



Reformulate our quality metrics:

- **Completeness:**

$$\frac{|\{v \in V \cup \Delta V \mid c_v + \Delta c_v = 1\}|}{|V \cup \Delta V|}$$

- **MutualInformation:**

$$\sum_{v_1} \sum_{v_2} \frac{c_{v_1 v_2} + \Delta c_{v_1 v_2}}{N + \Delta N} \log \frac{c_{v_1 v_2} + \Delta c_{v_1 v_2}}{(c_{v_1} + \Delta c_{v_1})(c_{v_2} + \Delta c_{v_2})}$$

```
val completeness = Completeness("origin")

// Compute state of the changed partition
val newStateToday =
  completeness.computeStateFrom(newPartitionToday)

// Load states of non-changed partitions
val (stateSunday, stateMonday) = loadPreviousStates(...)


// Sum of the states of the individual partitions
val newTableState = stateSunday + stateMonday + newStateToday

// Compute the completeness of 'origin' in the whole table from
state
val newTableCompleteness = completeness.computeMetricFrom(newTableState)
```

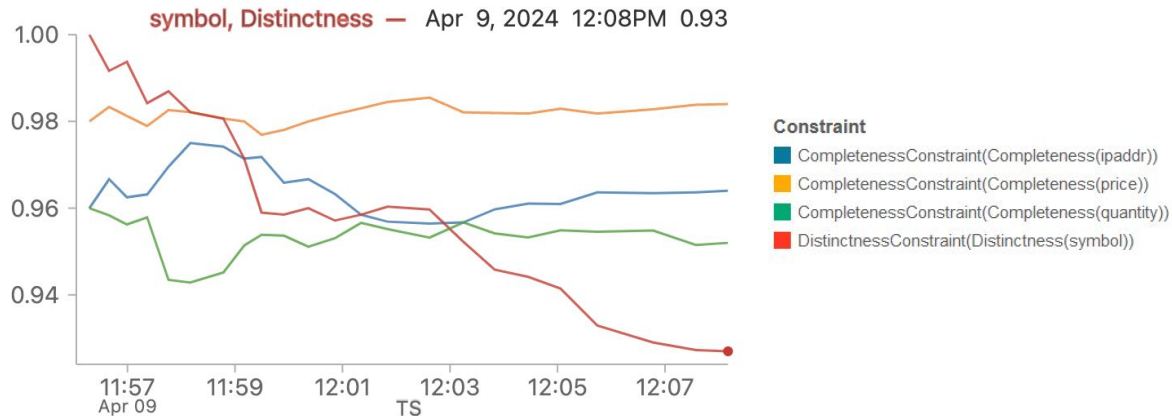
Appendix - 3 Example for an incremental update of the entropy of a column

4.1 Incremental Computation

In the following, we detail how to make our system's analyzers 'state-aware' to enable them to conduct incremental computations. A corresponding base class in Scala is shown in Listing 3, where M denotes the type of metric to compute and S denotes the type of state required. Persistence and retrieval of the state are handled outside of the implementation by a so-called `StateProvider`. The method `initialState` produces an initial empty state, `apply` produces a state and the corresponding metric for an initial dataset, and `update` consumes the current state and a delta dataset, and produces the updated state, as well as the corresponding metrics, both for the dataset as a whole and for the delta, in the form of a tuple (S, M, M) . Furthermore, the method `applyOrUpdateFromPersistedState` executes the incremental computation and takes care of managing the involved states using `StateProviders`.



```
1 trait IncrementalAnalyzer[M, S]
2   extends Analyzer[M] {
3
4   def initialState(initialData: DataFrame): S
5
6   def update(
7     state: S,
8     delta: DataFrame): (S, M, M)
9
10  def updateFromPersistedState(
11    stateProvider: Option[StateProvider],
12    nextStateProvider: StateProvider,
13    delta: DataFrame): (M, M)
14  }
15
16 trait StateProvider {
17
18  def persistState[S](
19    state: S,
20    analyzer: IncrementalAnalyzer[M, S])
21
22  def loadState[S](
23    analyzer: IncrementalAnalyzer[M, S]): S
24 }
```



Appendix - 2 Budget Control (AWS sucks)

CloudFormation > Stacks > databricks-workspace-stack-71f6f

Stacks (3)

Filter by stack name

Filter status: Deleted View nested


Stacks

- databricks-workspace-stack-71f6f
 - 2024-03-07 13:42:43 UTC-0800
 - DELETE_COMPLETE

Events (77)

Search events

Timestamp	Logical ID	Status
2024-03-07 13:46:38 UTC-0800	createWorkspace	CREATE_COMPLETE
2024-03-07 13:46:38 UTC-0800	createWorkspace	CREATE_IN_PROGRESS
2024-03-07 13:43:58 UTC-0800	createWorkspace	CREATE_IN_PROGRESS
2024-03-07 13:43:56 UTC-0800	createCredentials	CREATE_COMPLETE
2024-03-07 13:43:56 UTC-0800	createCredentials	CREATE_IN_PROGRESS



Total cost (historical and forecasted)

\$63.98

Average historical cost per month

\$21.30

Average forecasted cost per month

\$0.08

Costs (\$)



Appendix - 4 Delta Lake

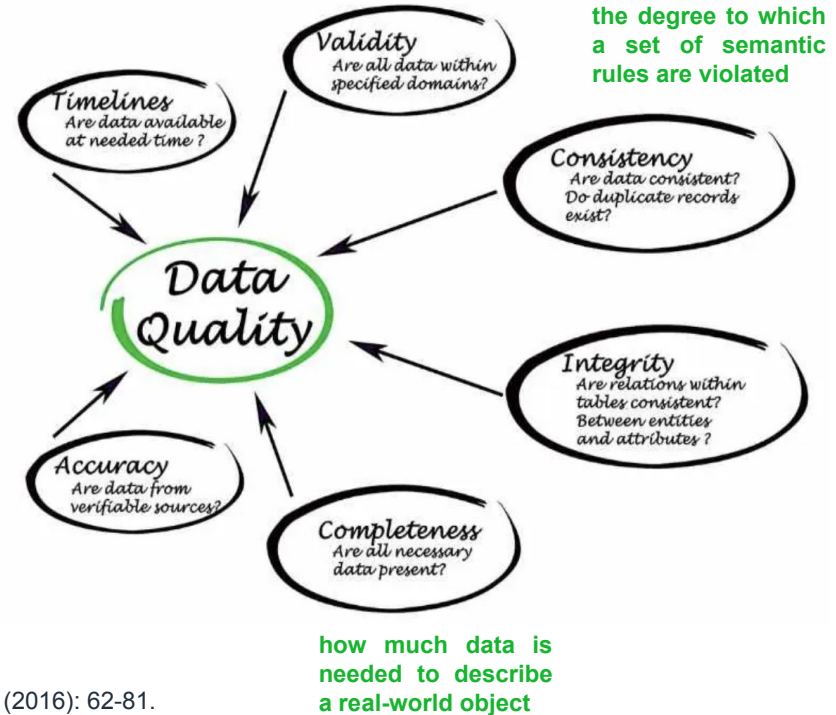
1. An optimized storage layer that provides the foundation for tables on Databricks.
2. It extends Parquet data files with a file-based transaction log for ACID transactions and scalable metadata handling.
3. Delta Lake is fully compatible with Apache Spark APIs, and was developed for tight integration with **Structured Streaming**, allowing users to easily use a single copy of data for both batch and streaming operations and providing incremental processing at scale.



Appendix - 5 What Is the Common Dimensions of Data Quality?

Challenges and derived requirements for various data domain:

Data Domain	Data Challenges	Derived Requirements
Company data	Duplicates	Consistent Representation
	Irrelevant data	Relevance
	Missing reference data	Value-Added
People data	Incomplete data	Accuracy
	Outdated data	Value-Added
	Root cause analysis	Timeliness
	Tracking issues	Relevance
Service/Asset data	Missing data	Accessibility
	Incomplete data	Value-added
	Incorrect values	Completeness
	Duplicates	Interpretability
Supply Chain data	Design errors	Completeness
	Duplicates	Timeliness
	Master data issues	Accuracy



[2] Silvola, Risto et al. "Data quality assessment and improvement." Int. J. Bus. Inf. Syst. 22 (2016): 62-81.